

Design and Implementation of Programming Languages

Project Specification

Fall 2019

Set-up: For this assignment, edit a copy of `project.rkt`, which is attached. In particular, replace occurrences of "CHANGE" to complete the problems. Do not use expressions with side effects (`set!`, `set-mcar!`, etc.) in your code.

Overview: This project has to do with NUMEX (**N**umber-**E**xpression Programming Language). NUMEX programs are written directly in Racket by using the constructors defined by the structs defined at the beginning of `project.rkt` (Note: you must define missing ones). Here is the definition of NUMEX's syntax:

- If s is a Racket string, then `(var s)` is a NUMEX expression (variables).
- If n is a Racket integer, then `(num n)` is a NUMEX expression (number constants).
- If b is a Racket boolean, then `(bool b)` is a NUMEX expression (boolean constants).
- If e_1 and e_2 are NUMEX expressions, then `(plus e_1 e_2)` is a NUMEX expression (addition).
- If e_1 and e_2 are NUMEX expressions, then `(minus e_1 e_2)` is a NUMEX expression (subtraction).
- If e_1 and e_2 are NUMEX expressions, then `(mult e_1 e_2)` is a NUMEX expression (multiplication).
- If e_1 and e_2 are NUMEX expressions, then `(div e_1 e_2)` is a NUMEX expression (division).
- If e_1 is a NUMEX expression, then `(neg e_1)` is a NUMEX expression (negation).
- If e_1 and e_2 are NUMEX expressions, then `(andalso e_1 e_2)` is a NUMEX expression (logical conjunction).
- If e_1 and e_2 are NUMEX expressions, then `(orelse e_1 e_2)` is a NUMEX expression (logical disjunction).
- If e_1 , e_2 , and e_3 are NUMEX expressions, then `(cnd e_1 e_2 e_3)` is a NUMEX expression. It is a condition where the result is e_2 if e_1 is true, else the result is e_3 . Only one of e_2 and e_3 is evaluated.
- If e_1 and e_2 are NUMEX expressions, then `(iseq e_1 e_2)` is a NUMEX expression (comparison).
- If e_1 , e_2 , and e_3 are NUMEX expressions, then `(ifnzero e_1 e_2 e_3)` is a NUMEX expression. It is a condition where the result is e_2 if e_1 is not zero, else the result is e_3 . Only one of e_2 and e_3 is evaluated.

- If e_1 , e_2 , e_3 , and e_4 are NUMEX expressions, then $(\text{ifleq } e_1 \ e_2 \ e_3 \ e_4)$ is a NUMEX expression. It is a conditional where the result is e_4 if e_1 is strictly greater than e_2 , else the result is e_3 . Only one of e_3 and e_4 is evaluated.
- If s_1 and s_2 are Racket strings and e is a NUMEX expression, then $(\text{lam } s_1 \ s_2 \ e)$ is a NUMEX expression (a function). In e , s_1 is bound to the function itself (for recursion) and s_2 is bound to the only argument. Also, $(\text{lam null } s_2 \ e)$ is allowed for anonymous nonrecursive functions.
- If e_1 and e_2 are NUMEX expressions, then $(\text{apply } e_1 \ e_2)$ is a NUMEX expression (function application).
- If s is a Racket string, and e_1 and e_2 are NUMEX expressions, then $(\text{with } s \ e_1 \ e_2)$ is a NUMEX expression (a let expression where the value of e_1 is bound to s in e_2).
- If e_1 and e_2 are NUMEX expressions, then $(\text{apair } e_1 \ e_2)$ is a NUMEX expression (pair constructor).
- If e_1 is a NUMEX expression, then $(\text{1st } e_1)$ is a NUMEX expression (the first part of a pair).
- If e_1 is a NUMEX expression, then $(\text{2nd } e_1)$ is a NUMEX expression (the second part of a pair).
- (munit) is a NUMEX expression (holding no data, much like $()$ in ML or null in Racket). Notice (munit) is a NUMEX expression, but `munit` is not.
- If e_1 is a NUMEX expression, then $(\text{ismunit } e_1)$ is a NUMEX expression (testing for (munit)).
- $(\text{closure } env \ f)$ is a NUMEX value where f is a NUMEX function and env is an environment that maps variables to values. Closures do not appear in programs; they result from evaluating functions.
- If s_1 is a Racket string and s_2 is a Racket string and e_1 is a NUMEX expression and e_2 is a NUMEX expression and e_3 is a NUMEX expression, then $(\text{letrec } s_1 \ e_1 \ s_2 \ e_2 \ e_3)$ is a NUMEX expression (a letrec expression for recursive definitions where the value of e_1 is bound to s_1 and the value of e_2 is bound to s_2 in the e_3).
- If s is a Racket string and e is a NUMEX expression, then $(\text{key } s \ e)$ is a NUMEX expression (key constructor).
- If k is a NUMEX key and m is a NUMEX munit, then $(\text{record } k \ m)$ is a NUMEX expression (record constructor).
- If k is a NUMEX key and r is a NUMEX record, then $(\text{record } k \ r)$ is a NUMEX expression (record constructor).
- If s is a Racket string and r is a NUMEX record, then $(\text{value } s \ r)$ is a NUMEX expression (value of string in record).

A NUMEX *value* is a NUMEX number constant, a NUMEX boolean constant, a NUMEX closure, a NUMEX munit, or a NUMEX pair of NUMEX values. Similar to Racket, we can build list values out of nested pair values that end with a NUMEX munit. Such a NUMEX value is called a NUMEX list.

You should *not* assume NUMEX programs are syntactically correct (e.g., things like (num "hi") or (num (num 37)) must be handled). And do *not* assume NUMEX programs are free of type errors like (plus (munit) (num 7)), (1st (num 7)) or (div (bool #t) (num 2)).

Instructions: Upload your modified `project.rkt` and `projectTest.rkt` through the Moodle website.

Problems:

1. Warm-Up

- (a) Write a Racket function `racketlist->numexlist` that takes a Racket list, which may even be a list of NUMEX values, and produces a NUMEX list with the same elements in the same order.
- (b) Write a Racket function `numexlist->racketlist` that takes a NUMEX list and produces a Racket list with the same elements in the same order.

2. Implementing NUMEX

Write an interpreter for NUMEX. It should be a Racket function `eval-exp` that takes a NUMEX expression `e` and either returns the NUMEX value that `e` evaluates to under the empty environment or calls Racket's error if evaluation encounters a run-time NUMEX type error or unbound NUMEX variables.

An NUMEX expression is evaluated in an environment (for evaluating variables, as usual). In your interpreter, use a Racket list of Racket pairs to represent this environment (which is initially empty) so that you can use the `envlookup` function, after completing it. Here is a description of the semantics of NUMEX expressions:

- All values (including closures) evaluate to themselves. For example, `(eval-exp (num 17))` would return `(num 17)`, not `17`.
- A variable evaluates to the value associated with it in the given environment.
- An arithmetic operation (addition, subtraction, multiplication, and division) evaluates to the result of what its operands evaluate to. Note: the operands must be numbers.
- A logical operation (`andalso` and `orelse`) evaluates to the result of what its operands evaluate to. Note: short-circuit evaluations are desired, and the operands must be booleans.

- A negation (`neg e`) evaluates to the opposite (negation) of what e evaluates to. Note: e can be a number or a boolean.
- For `(cnd e_1 e_2 e_3)`, the expression e_1 first evaluates to a boolean value. If the resulting value is `(bool #t)`, the whole expression evaluates to what e_2 evaluates to. The expression evaluates to the value of e_3 otherwise.
- The evaluation of `(iseq e_1 e_2)` involves the evaluation of e_1 and e_2 . The resulting value is `(bool #t)` if the value of e_1 equals to the value of e_2 . Otherwise, the expression evaluates to `(bool #f)`. Note: e_1 and e_2 can be numbers/booleans.
- For `(ifnzero e_1 e_2 e_3)`, the expression e_1 first evaluates to a value. If the resulting value is not zero, the whole expression evaluates to what e_2 evaluates to. The expression evaluates to the value of e_3 otherwise.
- The evaluation of `(ifleq e_1 e_2 e_3 e_4)` involves the evaluation of e_1 and e_2 . If the value of e_1 is strictly greater than the value of e_3 , then it evaluates to the value of e_4 . Otherwise, e_3 must be evaluated.
- For `(with s e_1 e_2)`, the expression e_2 evaluates to a value in an environment extended to map the name s to the evaluated value of e_1 .
- Functions are lexically scoped in the sense that a function evaluates to a closure holding the function and the current environment.
- For `(apply e_1 e_2)`, the expression e_1 first evaluates to a value. If the resulting value is not a closure, an error should be arisen. Otherwise, it evaluates the closure's function's body in the closure's environment extended to map the function's name to the closure (unless the name field is null) and the function's argument to the result of the evaluation of e_2 .
- The `(apair e_1 e_2)` construct makes a (new) pair holding the results of the evaluations of e_1 and e_2 .
- If the result of evaluating e_1 in `(1st e_1)` is an `apair`, then the first part is returned. Otherwise, it returns an error. Similarly, the evaluation of `(2nd e_1)` is the second part of the given pair.
- For `(ismunit e_1)`, the expression e_1 first evaluates to a value. If the resulting value is a `munit` expression, then the result is the NUMEX value `(bool #t)`, else it is the NUMEX value `(bool #f)`.
- For `(letrec s_1 e_1 s_2 e_2 e_3)` the expression e_3 evaluates to a value in an environment extended to map the name s_1 to the evaluated value of e_1 and the name s_2 to the evaluated value of e_2 .
- If s is a Racket string and the result of evaluating e is a NUMEX expression, the `(key s e)` construct makes a key holding corresponding value of s which is e . Otherwise, it returns an error.

- If the result of evaluating k is a key and the result of evaluating m is a munit, the `(record k m)` construct makes a record holding the result of the evaluation of k and the result of the evaluation of m . Otherwise, it returns an error.
- If the result of evaluating k is a key and the result of evaluating r is a record, the `(record k r)` construct makes a record holding the results of the evaluation of k and the evaluation of r . Otherwise, it returns an error.
- If the result of evaluating s is a string and the result of evaluating r is a record, the `(value s r)` returns corresponding value of s in r . If s does not exist in r , the `(value s r)` returns (munit). Otherwise, it returns an error.

3. Extending the Language

NUMEX is a small language, but we can write Racket functions that act like NUMEX macros so that users of these functions feel like NUMEX is larger. The Racket functions produce NUMEX expressions that could then be put inside larger NUMEX expressions or passed to `eval-exp`. In implementing these Racket functions, do not use closure (which is used only internally in `eval-exp`). Also do not use `eval-exp` (we are creating a program, not running it).

- Write a Racket function `ifmunit` that takes three NUMEX expressions e_1 , e_2 , and e_3 . It returns a NUMEX expression that first evaluates e_1 . If the resulting value is NUMEX's munit, then it evaluates e_2 and that is the overall result. Otherwise, e_3 must be evaluated.
- Write a Racket function `with*` that takes a Racket list of Racket pairs `'((s_1 . e_1) ... (s_i . e_i) ... (s_n . e_n))` and a final NUMEX expression e_{n+1} . In each pair, assume s_i is a Racket string and e_i is a NUMEX expression. `with*` returns a NUMEX expression whose value is e_{n+1} evaluated in an environment where each s_i is a variable bound to the result of evaluating the corresponding e_i for $1 \leq i \leq n$. The bindings are done sequentially, so that each e_i is evaluated in an environment where s_1 through s_{i-1} have been previously bound to the values e_1 through e_{i-1} .
- Write a Racket function `ifneq` that takes four NUMEX expressions e_1 , e_2 , e_3 , and e_4 and returns a NUMEX expression that acts like `ifleq` except e_3 is evaluated if and only if e_1 and e_2 are not equal numbers/booleans. Otherwise, the whole expression evaluates to what e_4 evaluates to. Assume none of the arguments to `ifneq` use the NUMEX variables `_x` or `_y`. Use this assumption so that when an expression returned from `ifneq` is evaluated, e_1 and e_2 are evaluated exactly once each.

4. Using the Language

We can write NUMEX expressions directly in Racket using the constructors for the structs and (for convenience) the functions we wrote in the previous problem.

- (a) Bind to the Racket variable `numex-filter` a NUMEX function that acts like `map` (as we use in ML). Your function should be curried: it should take a NUMEX function and return a NUMEX function that takes a NUMEX list and applies the function to every element of the list returning a new NUMEX list with all the elements for which the function returns a number other than zero (causing an error if the function returns a non-number). Recall a NUMEX list is `munit` or a pair where the second component is a NUMEX list.
- (b) Bind to the Racket variable `numex-all-gt` a NUMEX function that takes a NUMEX number i and returns a NUMEX function that takes a NUMEX list of NUMEX numbers and returns a new NUMEX list of NUMEX numbers containing the elements of the input list (in order) that are greater than i (hint: Use `numex-filter`).

5. Challenging Problem

Write a second version of `eval-exp` (bound to `eval-exp-c`) that builds closures with smaller environments: When building a closure, it uses an environment that is like the current environment but holds only variables that are free variables in the function part of the closure. (A free variable is a variable that appears in the function without being under some shadowing binding for the same variable.)

Avoid computing a function's free variables more than once. Do this by writing a function `compute-free-vars` that takes an expression and returns a different expression that uses `fun-challenge` everywhere in place of `fun`. The new struct `fun-challenge` (provided to you; do not change it) has a field `freevars` to store exactly the set of free variables for the function. Store this set as a Racket set of Racket strings. (Sets are predefined in Racket's standard library; consult the documentation for useful functions such as `set`, `set-add`, `set-member?`, `set-remove`, `set-union`, and any other functions you wish.)

You must have a top-level function `compute-free-vars` that works as just described — storing the free variables of each function in the `freevars` field — so the grader can test it directly. Then write a new “main part” of the interpreter that expects the sort of NUMEX expression that `compute-free-vars` returns. The case for function definitions is the interesting one.